

# project2

David Moran

2024-11-14

```
files <- list.files(path = "/Users/david_m123/Documents/gps", pattern = "*.geojson", full.names = TRUE)

all_features <- list()

for (file in files) {
  gps_data <- fromJSON(file, flatten = TRUE)
  features <- gps_data$features
  all_features <- append(all_features, list(features))
}

combined_features <- do.call(rbind,lapply(all_features,as.data.frame))

coordinates <- combined_features$geometry.coordinates
timestamps <- combined_features$properties.time
altitude <- combined_features$properties.altitude
accuracy <- combined_features$properties.accuracy
speed <- combined_features$properties.speed
bearings <- combined_features$properties.bearing

gps_df <- data.frame(
  longitude = sapply(coordinates, function(x) x[1]),
  latitude = sapply(coordinates, function(x) x[2]),
  timestamp = as.POSIXct(timestamps, format = "%Y-%m-%dT%H:%M:%OSZ", tz = "UTC"),
  altitude = altitude,
  speed = speed,
  bearings = bearings,
  stringsAsFactors = FALSE
)

head(gps_df)
```

```
##   longitude latitude      timestamp altitude speed bearings
## 1 -114.0003 46.88678 2020-08-18 17:50:42 971.2080  1.06   140.2
## 2 -114.0005 46.88632 2020-08-18 17:50:40 959.7000    NA     NA
## 3 -113.9999 46.88604 2020-08-18 17:51:49 972.2537  1.81    99.5
```

```
## 4 -113.9992 46.88620 2020-08-18 17:52:24 946.5977 3.41 93.1
## 5 -113.9978 46.88587 2020-08-18 17:53:12 979.1237 0.00 NA
## 6 -113.9972 46.88593 2020-08-18 17:53:11 960.9000 NA NA
```

```
gps_df_unique <- gps_df[!duplicated(gps_df[, c("longitude", "latitude")]), ]
num_removed <- nrow(gps_df) - nrow(gps_df_unique)
num_removed
```

```
## [1] 514
```

```
head(gps_df_unique)
```

```
##   longitude latitude      timestamp altitude speed bearings
## 1 -114.0003 46.88678 2020-08-18 17:50:42 971.2080 1.06 140.2
## 2 -114.0005 46.88632 2020-08-18 17:50:40 959.7000 NA NA
## 3 -113.9999 46.88604 2020-08-18 17:51:49 972.2537 1.81 99.5
## 4 -113.9992 46.88620 2020-08-18 17:52:24 946.5977 3.41 93.1
## 5 -113.9978 46.88587 2020-08-18 17:53:12 979.1237 0.00 NA
## 6 -113.9972 46.88593 2020-08-18 17:53:11 960.9000 NA NA
```

```
library(sf)
```

```
convert_to_utm <- function(gps_df_unique, zone = 12) {
  wgs84 <- st_crs(4326)

  utm_crs <- paste0("+proj=utm +zone=", zone, " +datum=WGS84")

  df_sf <- st_as_sf(gps_df_unique, coords = c("longitude", "latitude"), crs = wgs84)

  df_utm <- st_transform(df_sf, crs = utm_crs)

  coords <- st_coordinates(df_utm)
  gps_df_unique$easting <- coords[, 1]
  gps_df_unique$northing <- coords[, 2]
  return(gps_df_unique)
}
```

```
gps_data_utm <- convert_to_utm(gps_df_unique)
```

```
head(gps_data_utm)
```

```
##   longitude latitude      timestamp altitude speed bearings easting
## 1 -114.0003 46.88678 2020-08-18 17:50:42 971.2080 1.06 140.2 271422.2
## 2 -114.0005 46.88632 2020-08-18 17:50:40 959.7000 NA NA 271407.1
## 3 -113.9999 46.88604 2020-08-18 17:51:49 972.2537 1.81 99.5 271456.3
```

```
## 4 -113.9992 46.88620 2020-08-18 17:52:24 946.5977 3.41 93.1 271504.3
## 5 -113.9978 46.88587 2020-08-18 17:53:12 979.1237 0.00 NA 271614.3
## 6 -113.9972 46.88593 2020-08-18 17:53:11 960.9000 NA NA 271656.7
## northing
## 1 5196954
## 2 5196903
## 3 5196870
## 4 5196886
## 5 5196845
## 6 5196850
```

```
utm_data <- gps_data_utm[, c("timestamp", "easting", "northing", "altitude", "speed", "bearings")]

missing_values <- sapply(utm_data, function(x) sum(is.na(x)))

print(missing_values)
```

```
## timestamp easting northing altitude speed bearings
##          0         0         0        33    1664    1855
```

### *# Step 1: Data Exploration and Preprocessing*

#### *# 1.1 Visualize data to identify typical routes and anomalies*

#### *# 1.2 Handle Missing Values*

*# We will interpolate the speed and bearings columns since these values are continuous and can be estimated*

```
utm_data$speed <- zoo::na.approx(utm_data$speed, method = "linear")
utm_data$bearings <- zoo::na.approx(utm_data$bearings, method = "linear")
```

```
utm_data_cleaned <- na.omit(utm_data, cols = "altitude")
```

#### *# 1.3 Mark instances of missing or inconsistent GPS data*

*# For data consistency, we will create a new column that flags if the data had to be interpolated.*

```
utm_data_cleaned$interpolated <- apply(utm_data_cleaned, 1, function(row) {
  if (is.na(gps_data_utm[row["timestamp"], "speed"]) || is.na(gps_data_utm[row["timestamp"], "bearings"])) {
    return(1)
  } else {
    return(0)
  }
})
```

#### *# 1.4 Add Data Augmentation Features*

*# Adding additional features such as day of the week and time of day for analysis*

```
utm_data_cleaned$timestamp <- as.POSIXct(utm_data_cleaned$timestamp)
utm_data_cleaned$day_of_week <- weekdays(utm_data_cleaned$timestamp)
utm_data_cleaned$time_of_day <- sapply(format(utm_data_cleaned$timestamp, "%H"), function(x) {
```

```

if (x >= 5 & x < 12) {
  return("morning")
} else if (x >= 12 & x < 17) {
  return("afternoon")
} else if (x >= 17 & x < 21) {
  return("evening")
} else {
  return("night")
}
})

```

```
head(utm_data_cleaned)
```

```

##           timestamp easting northing altitude      speed bearings
## 1 2020-08-18 17:50:42 271422.2 5196954 971.2080 1.0600000 140.200
## 2 2020-08-18 17:50:40 271407.1 5196903 959.7000 1.4350000 119.850
## 3 2020-08-18 17:51:49 271456.3 5196870 972.2537 1.8100000  99.500
## 4 2020-08-18 17:52:24 271504.3 5196886 946.5977 3.4100000  93.100
## 5 2020-08-18 17:53:12 271614.3 5196845 979.1237 0.0000000 105.925
## 6 2020-08-18 17:53:11 271656.7 5196850 960.9000 0.1357143 118.750
##   interpolated day_of_week time_of_day
## 1             1    Tuesday    evening
## 2             1    Tuesday    evening
## 3             1    Tuesday    evening
## 4             1    Tuesday    evening
## 5             1    Tuesday    evening
## 6             1    Tuesday    evening

```

### *# 2.1 Calculate Basic Features: Speed, Bearing, Acceleration, and Distance*

```
utm_data_cleaned$acceleration <- c(NA, diff(utm_data_cleaned$speed) / as.numeric(diff(utm_data_cleaned$timestamp)))
```

```
utm_data_cleaned$distance <- c(NA, sqrt(diff(utm_data_cleaned$easting)^2 + diff(utm_data_cleaned$northing)^2))
```

### *# 2.2 Identify Transit Periods*

```
speed_threshold <- 0.5
```

```
utm_data_cleaned$in_transit <- ifelse(utm_data_cleaned$speed > speed_threshold, 1, 0)
```

### *# 2.3 Time-Based Feature Set*

```
# Identify peak transit hours
```

```
# Create a new column for hour of the day
```

```
utm_data_cleaned$hour <- as.integer(format(utm_data_cleaned$timestamp, "%H"))
```

```
peak_hours <- table(utm_data_cleaned$hour)
```

### *# 2.4 Extract Frequent Routes using Clustering*

```
# Clustering the UTM coordinates to find frequently traveled routes
```

```
coords <- as.matrix(utm_data_cleaned[, c("easting", "northing")])
db <- dbscan::dbscan(coords, eps = 50, minPts = 5)
```

```
utm_data_cleaned$route_cluster <- db$cluster
```

```
head(utm_data_cleaned)
```


```
##           timestamp easting northing altitude      speed bearings
## 1 2020-08-18 17:50:42 271422.2 5196954 971.2080 1.0600000 140.200
## 2 2020-08-18 17:50:40 271407.1 5196903 959.7000 1.4350000 119.850
## 3 2020-08-18 17:51:49 271456.3 5196870 972.2537 1.8100000  99.500
## 4 2020-08-18 17:52:24 271504.3 5196886 946.5977 3.4100000  93.100
## 5 2020-08-18 17:53:12 271614.3 5196845 979.1237 0.0000000 105.925
## 6 2020-08-18 17:53:11 271656.7 5196850 960.9000 0.1357143 118.750
##   interpolated day_of_week time_of_day acceleration distance in_transit hour
## 1             1    Tuesday    evening             NA         NA           1   17
## 2             1    Tuesday    evening -0.354777699  52.69876           1   17
## 3             1    Tuesday    evening  0.005510087  59.31356           1   17
## 4             1    Tuesday    evening  0.045714286  50.52923           1   17
## 5             1    Tuesday    evening -0.071041667 117.48222           0   17
## 6             1    Tuesday    evening -0.892857434  42.77727           0   17
##   route_cluster
## 1             1
## 2             1
## 3             1
## 4             1
## 5             1
## 6             1
```

```
#Visualizations for Feature Analysis
```

```
# Visualization 1: Histogram of Peak Transit Hours
```

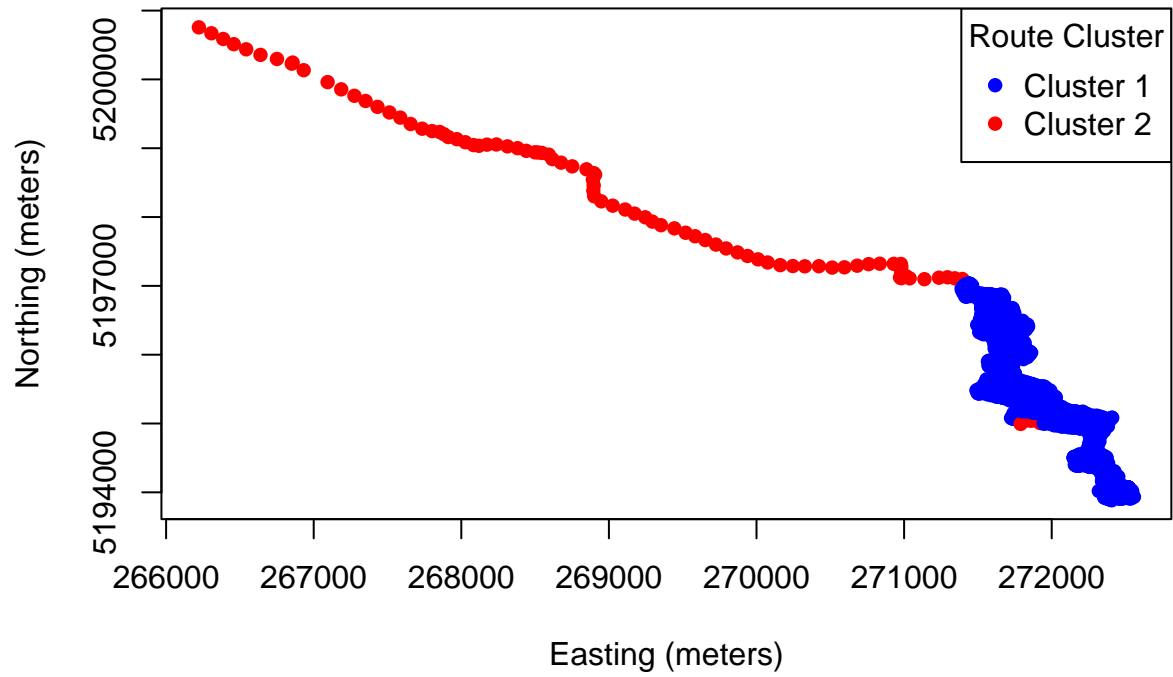
```
hist(utm_data_cleaned$hour, breaks = 24, main = "Histogram of User Transit Hours", xlab = "Hour of Day"
```

## Histogram of User Transit Hours



```
# Visualization 2: Scatter Plot of User Movement (Easting vs Northing with Cluster Labels)
plot(utm_data_cleaned$easting, utm_data_cleaned$northing, col = ifelse(utm_data_cleaned$route_cluster ==
legend("topright", legend = c("Cluster 1", "Cluster 2"), col = c("blue", "red"), pch = 16, title = "Rou
```

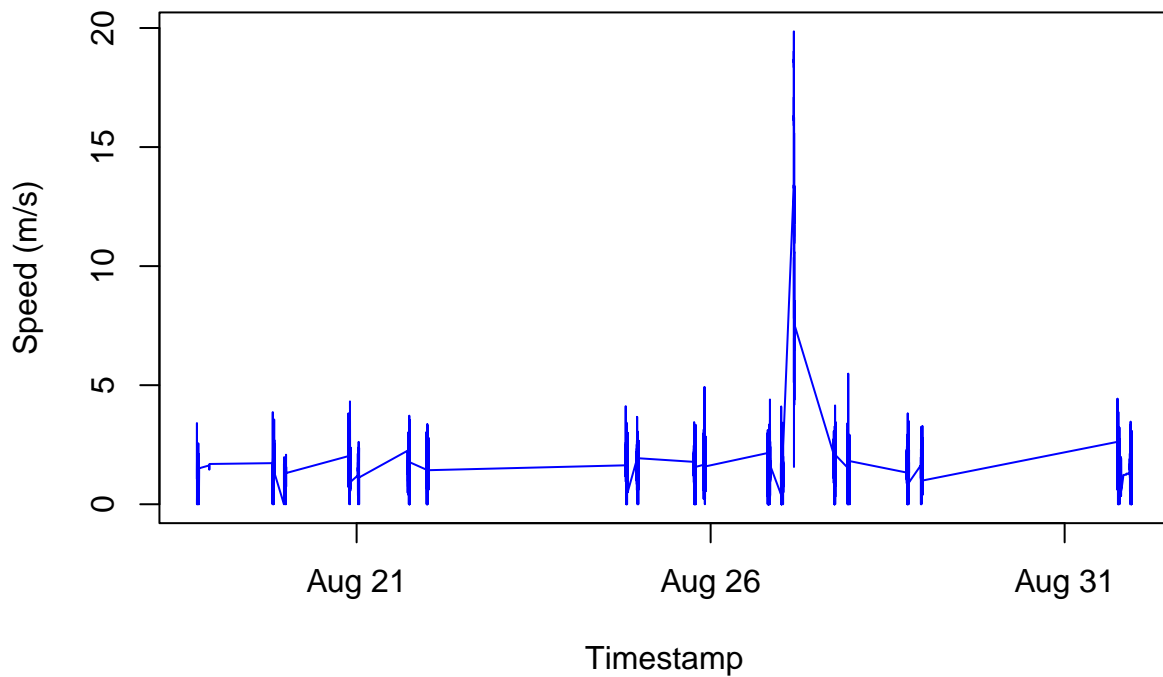
## User Movement Paths with Route Clustering



```
# Visualization 3: Line Plot of Speed over Time
```

```
plot(utm_data_cleaned$timestamp, utm_data_cleaned$speed, type = "l", col = "blue", xlab = "Timestamp", ylab = "Speed (m/s)",
```

## User Speed over Time



```
# Analyzing Identified Route Clusters
```

```
cluster_counts <- table(utm_data_cleaned$route_cluster)
print("Cluster Counts (Number of points per cluster):")
```

```
## [1] "Cluster Counts (Number of points per cluster):"
```

```
print(cluster_counts)
```

```
##
##      0      1
## 91 5219
```

```
# Calculate basic statistics for each route cluster (e.g., average speed, average distance traveled with...
route_cluster_analysis <- aggregate(cbind(speed, distance, acceleration) ~ route_cluster, data = utm_da
```

```
route_cluster_analysis <- aggregate(cbind(speed, distance, acceleration) ~ route_cluster, data = utm_da
```


```
route_cluster_analysis <- do.call(data.frame, route_cluster_analysis)
```



```
colnames(route_cluster_analysis) <- c("route_cluster", "avg_speed", "speed_std_dev", "avg_distance", "distance_std_dev")
print(route_cluster_analysis)
```

```
##   route_cluster avg_speed speed_std_dev avg_distance distance_std_dev
## 1             0 11.128702    4.5494347   153.81471    664.93758
## 2             1  1.491226    0.6265929    25.93676     40.20385
##   avg_acceleration acceleration_std_dev
## 1      0.001975239          0.3699626
## 2     -0.025373930          1.6308716
```

```
# Visualization: Average Speed by Route Cluster
barplot(route_cluster_analysis$avg_speed, names.arg = route_cluster_analysis$route_cluster, col = "skyblue")
```



```
# Filter out noise points (cluster -1)
filtered_data <- subset(utm_data_cleaned, route_cluster != -1)

# Display the filtered dataset to the user
head(filtered_data)
```

```
##           timestamp easting northing altitude      speed bearings
## 1 2020-08-18 17:50:42 271422.2  5196954 971.2080 1.0600000   140.200
## 2 2020-08-18 17:50:40 271407.1  5196903 959.7000 1.4350000   119.850
## 3 2020-08-18 17:51:49 271456.3  5196870 972.2537 1.8100000    99.500
```


```
## 4 2020-08-18 17:52:24 271504.3 5196886 946.5977 3.4100000 93.100
## 5 2020-08-18 17:53:12 271614.3 5196845 979.1237 0.0000000 105.925
## 6 2020-08-18 17:53:11 271656.7 5196850 960.9000 0.1357143 118.750
##   interpolated day_of_week time_of_day acceleration distance in_transit hour
## 1           1    Tuesday    evening           NA         NA           1    17
## 2           1    Tuesday    evening -0.354777699  52.69876           1    17
## 3           1    Tuesday    evening  0.005510087  59.31356           1    17
## 4           1    Tuesday    evening  0.045714286  50.52923           1    17
## 5           1    Tuesday    evening -0.071041667 117.48222           0    17
## 6           1    Tuesday    evening -0.892857434  42.77727           0    17
##   route_cluster
## 1             1
## 2             1
## 3             1
## 4             1
## 5             1
## 6             1
```

```
# Visualize Movement Paths without Noise Cluster
```

```
# Scatter Plot of User Movement (Easting vs Northing with Filtered Cluster Labels)
```

```
plot(filtered_data$easting, filtered_data$northing, col = filtered_data$route_cluster, pch = 16, xlab =
legend("topright", legend = unique(filtered_data$route_cluster), col = unique(filtered_data$route_cluster))
```

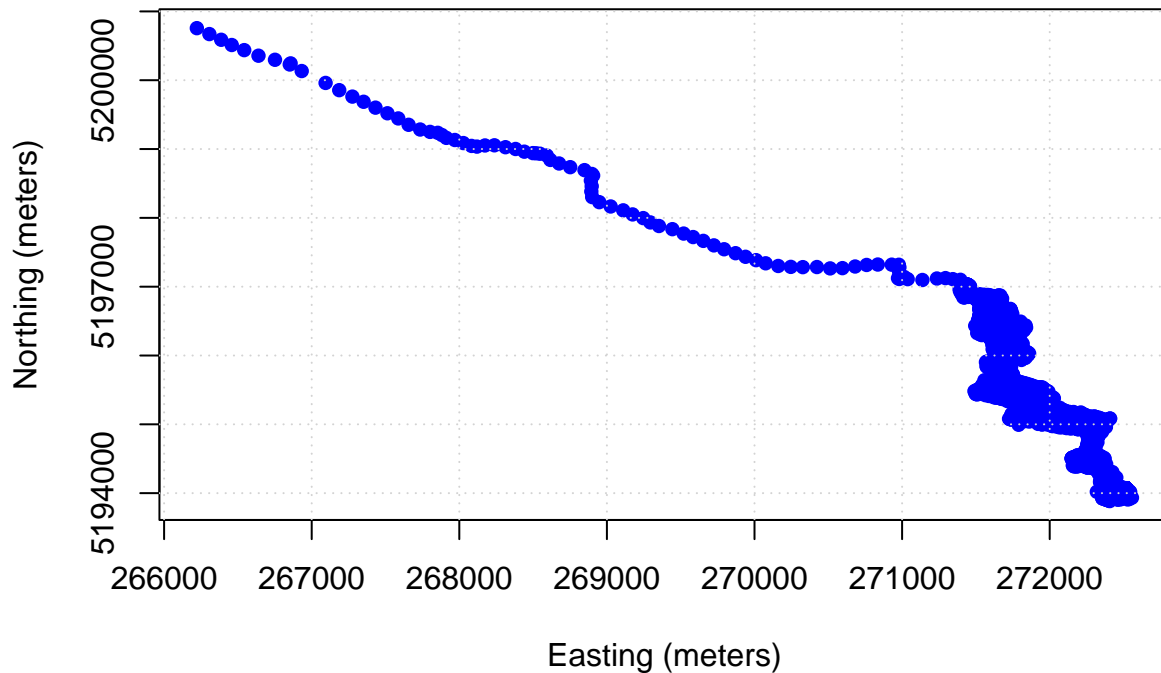
## User Movement Paths without Noise Cluster



```
# Scatter Plot of User Movement without Noise Cluster (Clearer Visualization)
```

```
plot(filtered_data$easting, filtered_data$northing, col = "blue", pch = 16, xlab = "Easting (meters)", ylab = "Northing (meters)", grid())
```

## User Movement Paths without Noise Cluster



```
# Step 3: Feature Extraction for Pattern Prediction
```

```
# Step 3.1: Extract Features for Predicting Transit Periods
```

```
# Create a new column to identify "trip segments" - periods when the user is in transit
```

```
filtered_data$trip_segment <- cumsum(c(1, diff(filtered_data$in_transit)) != 0)
filtered_data$trip_segment <- ifelse(filtered_data$in_transit == 1, filtered_data$trip_segment, NA)
filtered_data$trip_segment <- zoo::na.locf(filtered_data$trip_segment, na.rm = FALSE)
```

```
# Step 3.2: Calculate Summary Statistics for Each Trip Segment
```

```
# This includes average speed, distance, acceleration, and the length of the trip
```

```
trip_summary <- filtered_data %>%
  group_by(trip_segment) %>%
  summarise(
    avg_speed = mean(speed, na.rm = TRUE),
    max_speed = max(speed, na.rm = TRUE),
    speed_std_dev = sd(speed, na.rm = TRUE),
    total_distance = sum(distance, na.rm = TRUE),
    avg_acceleration = mean(acceleration, na.rm = TRUE),
    start_time = min(timestamp, na.rm = TRUE),
    end_time = max(timestamp, na.rm = TRUE)
```

```

) %>%
ungroup()

trip_summary$duration_seconds <- as.numeric(difftime(trip_summary$end_time, trip_summary$start_time, units = "secs"))

trip_summary$duration_seconds <- as.numeric(difftime(as.POSIXct(trip_summary$end_time, origin = '1970-01-01'), as.POSIXct(trip_summary$start_time, origin = '1970-01-01'), units = "secs"))

head(trip_summary)

```

```

## # A tibble: 6 x 9
##   trip_segment avg_speed max_speed speed~1 total~2 avg_ac~3 start_time
##   <int>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl> <dtm>
## 1         1      1.07      3.41     1.15     473.    -1.80e-1 2020-08-18 17:50:40
## 2         3      0.532     0.95     0.313     533.    -2.32e-2 2020-08-18 17:54:46
## 3         5      0.843     1.62     0.457    1515.    -7.04e-4 2020-08-18 18:00:36
## 4         7      1.35     2.56     0.732    1465.    -2.63e-2 2020-08-18 18:15:14
## 5         9      0.905     1.67     0.615     417.     6.89e-2 2020-08-18 18:29:52
## 6        11      1.32     1.76     0.482    3920.    -1.00e-2 2020-08-18 18:34:35
## # ... with 2 more variables: end_time <dtm>, duration_seconds <dbl>, and
## #   abbreviated variable names 1: speed_std_dev, 2: total_distance,
## #   3: avg_acceleration

```

*# Step 3.3: Time Prediction Model Using K-Means Clustering in R*

```

library(dplyr)
library(lubridate)

```

```

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

```

```

library(stats)

```

```

# Extract the hour and minute from the timestamp for time clustering
filtered_data <- filtered_data %>%
  mutate(hour_minute = hour(timestamp) * 60 + minute(timestamp))

```

```

departure_times <- trip_summary %>%
  select(start_time) %>%
  mutate(hour_minute = hour(start_time) * 60 + minute(start_time))

```

```

# Use K-Means clustering to identify common departure time patterns (e.g., morning vs. evening departures)
set.seed(42)
kmeans_result <- kmeans(departure_times$hour_minute, centers = 3)

```

```
departure_times$cluster <- kmeans_result$cluster
```

```
# Determine the centroids (average departure times) for each cluster
```

```
departure_time_centroids <- kmeans_result$centers
```

```
centroids_converted <- data.frame(  
  Hour = floor(departure_time_centroids / 60),  
  Minute = round(departure_time_centroids %% 60)  
)
```

```
print("Predicted Common Departure Time Centroids")
```

```
## [1] "Predicted Common Departure Time Centroids"
```

```
print(centroids_converted)
```

```
##   Hour Minute  
## 1    22     55  
## 2     0     21  
## 3    18     55
```

```
print("Departure Times with Predicted Clusters")
```

```
## [1] "Departure Times with Predicted Clusters"
```

```
head(departure_times)
```

```
## # A tibble: 6 x 3  
##   start_time      hour_minute cluster  
##   <dtm>          <dbl>    <int>  
## 1 2020-08-18 17:50:40      1070      3  
## 2 2020-08-18 17:54:46      1074      3  
## 3 2020-08-18 18:00:36      1080      3  
## 4 2020-08-18 18:15:14      1095      3  
## 5 2020-08-18 18:29:52      1109      3  
## 6 2020-08-18 18:34:35      1114      3
```

```
# Step 3.4: Optimized Path and Location Prediction Model Using Linear Interpolation in R
```

```
library(dplyr)  
library(lubridate)
```

```
# Use departure time centroids and path predictions to generate synthetic week 3 data more efficiently  
week3_predictions <- list()
```

```
# Use departure time centroids and generate synthetic week 3 data using vectorized operations  
for (i in 1:nrow(centroids_converted)) {
```

```

start_hour <- centroids_converted$Hour[i]
start_minute <- centroids_converted$Minute[i]
departure_time <- make_datetime(year = 2020, month = 9, day = 1, hour = start_hour, min = start_minute)

# Predict the initial cluster and start location based on historical patterns
likely_cluster <- as.numeric(names(sort(table(filtered_data$route_cluster), decreasing = TRUE)[1]))
cluster_data <- filtered_data %>%
  filter(route_cluster == likely_cluster) %>%
  mutate(time_offset = as.numeric(difftime(timestamp, min(timestamp), units = "secs")))

cluster_data <- cluster_data %>%
  mutate(new_timestamp = departure_time + seconds(time_offset))

week3_predictions[[i]] <- cluster_data %>%
  select(new_timestamp, easting, northing, speed, acceleration, route_cluster)
}

week3_predictions_df <- bind_rows(week3_predictions) %>%
  rename(timestamp = new_timestamp)

print("Optimized Predicted GPS Data for Week 3 (Interpolation)")

```

```
## [1] "Optimized Predicted GPS Data for Week 3 (Interpolation)"
```

```
head(week3_predictions_df)
```

```
##           timestamp easting northing    speed acceleration route_cluster
## 1 2020-09-01 22:55:01 271422.2 5196954 1.0600000          NA           1
## 2 2020-09-01 22:55:00 271407.1 5196903 1.4350000 -0.354777699           1
## 3 2020-09-01 22:56:08 271456.3 5196870 1.8100000 0.005510087           1
## 4 2020-09-01 22:56:43 271504.3 5196886 3.4100000 0.045714286           1
## 5 2020-09-01 22:57:31 271614.3 5196845 0.0000000 -0.071041667           1
## 6 2020-09-01 22:57:30 271656.7 5196850 0.1357143 -0.892857434           1
```

```
# Step 3.5: Implementing the Tagging Algorithm for Week 3 in R
```

```

tagging_algorithm <- function(week3_data) {
  tagged_points <- list()

  current_trip <- NA_character_
  start_time <- NA

  for (i in 1:nrow(week3_data)) {
    row <- week3_data[i, ]

```

```

if (!is.na(row$route_cluster) && (is.na(current_trip) || !isTRUE(all.equal(row$route_cluster, current_trip)))) {

  current_trip <- row$route_cluster
  start_time <- row$timestamp
}

# Apply tagging rules:
if (!is.na(start_time) && !is.na(row$timestamp)) {
  time_since_start <- as.numeric(difftime(row$timestamp, start_time, units = "secs"))

  # Only tag if more than 5 minutes have passed since the start of transit
  if (time_since_start > 300) {
    # Avoid tagging if stationary for more than 2 minutes
    if (!is.na(row$speed) && row$speed > 0.5) { # Using a threshold of 0.5 m/s for stationary detection
      tagged_points <- append(tagged_points, list(row))
    }
  }
}

if (length(tagged_points) > 0) {
  tagged_points_df <- do.call(rbind, tagged_points) %>% as.data.frame()
} else {
  tagged_points_df <- data.frame(timestamp = as.POSIXct(character()),
                                easting = numeric(),
                                northing = numeric(),
                                speed = numeric(),
                                acceleration = numeric(),
                                route_cluster = integer())
}

return(tagged_points_df)
}

week3_tagged_df <- tagging_algorithm(week3_predictions_df)

print("Tagged GPS Data for Week 3")

```

```
## [1] "Tagged GPS Data for Week 3"
```

```
head(week3_tagged_df)
```

```
##           timestamp easting northing    speed acceleration route_cluster
## 11 2020-09-01 23:00:27 271680.7 5196573 0.8142857 0.002690448          1
## 12 2020-09-01 23:01:34 271645.1 5196499 0.9500000 0.002040295          1
## 13 2020-09-01 23:02:09 271638.4 5196465 0.8100000 -0.004000000          1
## 18 2020-09-01 23:04:55 271710.9 5196245 0.5900000 0.003787076          1
## 19 2020-09-01 23:05:37 271694.4 5196151 0.7866667 0.004652520          1
## 20 2020-09-01 23:05:58 271652.5 5196144 0.9833333 0.009458766          1
```

*# Step 3.6: Implementing a Manual Kalman Filter for Location Prediction in R*

```
library(Matrix)
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following object is masked from 'package:dlm':
```

```
##
```

```
##      bdiag
```

```
# Kalman Filter parameters
```

```
dt <- 1 # Time step (seconds)
```

```
A <- matrix(c(1, 0, dt, 0,  
             0, 1, 0, dt,  
             0, 0, 1, 0,  
             0, 0, 0, 1), nrow = 4, byrow = TRUE) # State transition matrix
```

```
H <- matrix(c(1, 0, 0, 0,  
             0, 1, 0, 0), nrow = 2, byrow = TRUE) # Observation matrix
```

```
Q <- diag(4) * 0.1 # Process noise covariance
```

```
R <- diag(2) * 5 # Measurement noise covariance
```

```
P <- diag(4) # Initial estimate error covariance
```

```
# Initialize state vector (easting, northing, velocity in easting, velocity in northing)
```

```
initial_position <- as.numeric(week3_predictions_df[1, c("easting", "northing")])
```

```
initial_velocity <- c(0, 0)
```

```
x <- c(initial_position, initial_velocity) # Initial state
```

```
kalman_predictions <- list()
```

```
# Run the Kalman Filter for each timestamp in week 3 predictions
```

```
for (idx in 1:nrow(week3_predictions_df)) {
```

```
  # Prediction Step
```

```
  x <- A %*% x
```

```
  P <- A %*% P %*% t(A) + Q
```

```
  # Update Step (only if we have observations)
```

```
  z <- as.numeric(week3_predictions_df[idx, c("easting", "northing")]) # Observed position
```

```
  y <- z - (H %*% x) # Measurement residual
```

```
  S <- H %*% P %*% t(H) + R # Residual covariance
```

```
  K <- P %*% t(H) %*% solve(S) # Kalman gain
```

```
  x <- x + (K %*% y) # Updated state estimate
```

```
  P <- (diag(4) - K %*% H) %*% P # Updated estimate covariance
```

```
kalman_predictions[[idx]] <- data.frame(  
  timestamp = week3_predictions_df$timestamp[idx],  
  easting = x[1],  
  northing = x[2],
```



```

    velocity_easting = x[3],
    velocity_northing = x[4],
    route_cluster = week3_predictions_df$route_cluster[idx]
  )
}

```

```
kalman_predictions_df <- do.call(rbind, kalman_predictions)
```

```
print("Kalman Filter Predicted GPS Data for Week 3")
```

```
## [1] "Kalman Filter Predicted GPS Data for Week 3"
```

```
head(kalman_predictions_df)
```

```
##           timestamp easting northing velocity_easting velocity_northing
## 1 2020-09-01 22:55:01 271422.2 5196954         0.000000         0.000000
## 2 2020-09-01 22:55:00 271415.5 5196931        -2.808831        -9.386817
## 3 2020-09-01 22:56:08 271434.3 5196896         4.584480        -18.168107
## 4 2020-09-01 22:56:43 271471.1 5196882        13.967181        -17.016069
## 5 2020-09-01 22:57:31 271545.9 5196856        30.098393        -19.537567
## 6 2020-09-01 22:57:30 271612.5 5196843        39.291690        -17.914544
##   route_cluster
## 1             1
## 2             1
## 3             1
## 4             1
## 5             1
## 6             1
```

*# Step 3.7: Apply Tagging Algorithm to Kalman Filter Predicted Data for Week 3*

```

tagging_algorithm_kalman <- function(week3_data) {
  library(dplyr)
  library(lubridate)

  tagged_points <- list()

  current_trip <- NA
  start_time <- NA

  for (idx in 1:nrow(week3_data)) {
    row <- week3_data[idx, ]

    # Check for a new trip segment
    if (is.na(current_trip) || row$route_cluster != current_trip) {
      # Start of a new trip
      current_trip <- row$route_cluster
    }
  }
}

```

```

    start_time <- row$timestamp
  }

  # Calculate time since the start of the trip
  time_since_start <- as.numeric(difftime(row$timestamp, start_time, units = "secs"))

  # Only tag if more than 5 minutes have passed since the start of transit
  if (time_since_start > 300) {
    # Avoid tagging if stationary for more than 2 minutes
    velocity_magnitude <- sqrt(row$velocity_easting^2 + row$velocity_northing^2)
    if (velocity_magnitude > 0.5) { # Threshold of 0.5 m/s for stationary detection
      tagged_points <- append(tagged_points, list(row))
    }
  }
}

tagged_points_df <- bind_rows(tagged_points)

return(tagged_points_df)
}

```

```

week3_kalman_tagged_df <- tagging_algorithm_kalman(kalman_predictions_df)

```

```

print("Tagged GPS Data for Week 3 (Kalman Filter):")

```

```

## [1] "Tagged GPS Data for Week 3 (Kalman Filter):"

```

```

head(week3_kalman_tagged_df)

```

```

##           timestamp easting northing velocity_easting velocity_northing
## 1 2020-09-01 23:00:27 271674.8 5196605      17.40833463      -42.73772
## 2 2020-09-01 23:01:34 271672.2 5196535      12.36990353      -49.46006
## 3 2020-09-01 23:02:09 271665.0 5196477       7.41431454      -51.68392
## 4 2020-09-01 23:02:08 271637.9 5196441      -1.30313344      -47.85330
## 5 2020-09-01 23:03:33 271641.4 5196386      -0.08275271      -49.61005
## 6 2020-09-01 23:03:32 271647.5 5196326       1.47576998      -52.09136
## route_cluster
## 1             1
## 2             1
## 3             1
## 4             1
## 5             1
## 6             1

```

```

# Step 4: Revised Implementation Using Generated Synthetic Week 3 Data as Initial Input

```

```

library(dplyr)

```

```

# Step 4: Iterate Over Each Travel Segment from the Synthetic Week 3 Data
week3_segments <- split(kalman_predictions_df, kalman_predictions_df$route_cluster) # Group data by route_cluster

kalman_predictions_final <- list()

for (segment_id in names(week3_segments)) {
  segment_data <- week3_segments[[segment_id]]

  initial_row <- segment_data[1, ]
  initial_timestamp <- initial_row$timestamp
  initial_easting <- initial_row$easting
  initial_northing <- initial_row$northing

  # Initialize Kalman Filter with Initial Location of Each Travel Segment
  initial_position <- c(initial_easting, initial_northing)
  initial_velocity <- c(0, 0) # Assume starting from rest for each segment
  x <- c(initial_position, initial_velocity) # Initial state

  P <- diag(4)

  # Run Kalman Filter for each point in the travel segment
  for (i in 1:nrow(segment_data)) {
    row <- segment_data[i, ]

    time_offset <- as.numeric(difftime(row$timestamp, min(segment_data$timestamp), units = "secs"))
    new_timestamp <- initial_timestamp + time_offset

    # Prediction Step
    x <- A %>% x
    P <- A %>% P %>% t(A) + Q

    # Update Step
    z <- c(row$easting, row$northing)
    y <- z - (H %>% x) # Measurement residual
    S <- H %>% P %>% t(H) + R # Residual covariance
    K <- P %>% t(H) %>% solve(S) # Kalman gain
    x <- x + (K %>% y) # Updated state estimate
    P <- (diag(4) - K %>% H) %>% P # Updated estimate covariance

    kalman_predictions_final <- append(kalman_predictions_final, list(data.frame(
      timestamp = new_timestamp,
      easting = x[1],
      northing = x[2],
      velocity_easting = x[3],
      velocity_northing = x[4],
      route_cluster = segment_id
    )))
  }
}

```

```

kalman_predictions_final_df <- do.call(rbind, kalman_predictions_final)

week3_kalman_tagged_final_df <- tagging_algorithm_kalman(kalman_predictions_final_df)

print("Final Revised Tagged GPS Data for Week 3 (Kalman Filter):")

## [1] "Final Revised Tagged GPS Data for Week 3 (Kalman Filter):"

head(week3_kalman_tagged_final_df)

##           timestamp easting northing velocity_easting velocity_northing
## 1 2020-09-02 21:34:28 271697.7 5196634      22.8248864      -39.58494
## 2 2020-09-02 21:35:35 271700.0 5196570      17.6527025      -45.94198
## 3 2020-09-02 21:36:10 271695.3 5196504      12.0005679      -50.92806
## 4 2020-09-02 21:36:09 271677.8 5196448       4.5517277      -52.25263
## 5 2020-09-02 21:37:34 271665.0 5196391       0.1568505      -53.28884
## 6 2020-09-02 21:37:33 271657.6 5196333      -1.7341048      -54.54422
##   route_cluster
## 1             1
## 2             1
## 3             1
## 4             1
## 5             1
## 6             1

library(dplyr)
library(sf)

utm_crs <- 32633
latlon_crs <- 4326

week3_kalman_tagged_final_df <- week3_kalman_tagged_final_df %>%
  rowwise() %>%
  mutate(

    geometry = st_sfc(st_point(c(easting, northing)), crs = utm_crs),

    geometry_latlon = st_transform(geometry, crs = latlon_crs),

    latitude = st_coordinates(geometry_latlon)[2],
    longitude = st_coordinates(geometry_latlon)[1]
  ) %>%
  select(-geometry, -geometry_latlon) %>%
  ungroup()

print("Updated DataFrame with Latitude and Longitude:")

```

```
## [1] "Updated DataFrame with Latitude and Longitude:"
```

```
print(head(week3_kalman_tagged_final_df))
```

```
## # A tibble: 6 x 8
##   timestamp          easting northing velocit~1 veloc~2 route~3 latit~4 longi~5
##   <dtm>              <dbl>   <dbl>   <dbl>   <dbl> <chr>      <dbl>   <dbl>
## 1 2020-09-02 21:34:28 271698. 5196634.    22.8    -39.6 1        46.9    12.0
## 2 2020-09-02 21:35:35 271700. 5196570.    17.7    -45.9 1        46.9    12.0
## 3 2020-09-02 21:36:10 271695. 5196504.    12.0    -50.9 1        46.9    12.0
## 4 2020-09-02 21:36:09 271678. 5196448.     4.55   -52.3 1        46.9    12.0
## 5 2020-09-02 21:37:34 271665. 5196391.     0.157  -53.3 1        46.9    12.0
## 6 2020-09-02 21:37:33 271658. 5196333.    -1.73   -54.5 1        46.9    12.0
## # ... with abbreviated variable names 1: velocity_easting,
## #   2: velocity_northing, 3: route_cluster, 4: latitude, 5: longitude
```

```
library(dplyr)
```

```
# Define the cutoff date as September 7, 2020, at 23:59:59
cutoff_date <- as.POSIXct("2020-09-07 23:59:59", tz = "UTC")
```

```
# Truncate the synthetic Week 3 data to only include data until September 7
kalman_predictions_final_truncated_df <- kalman_predictions_final_df %>%
  filter(timestamp <= cutoff_date)
```

```
# Truncate the generated tagged data to only include tags until September 7
week3_kalman_tagged_final_truncated_df <- week3_kalman_tagged_final_df %>%
  filter(timestamp <= cutoff_date)
```

```
print("Truncated Synthetic Week 3 Data (Kalman Filter Predictions):")
```

```
## [1] "Truncated Synthetic Week 3 Data (Kalman Filter Predictions):"
```

```
print(head(kalman_predictions_final_truncated_df))
```

```
##           timestamp  easting northing velocity_easting velocity_northing
## 1 2020-09-02 21:29:02 271422.2 5196954         0.000000         0.000000
## 2 2020-09-02 21:29:01 271419.2 5196944        -1.239034        -4.140721
## 3 2020-09-02 21:30:09 271426.1 5196918         1.529014       -11.490703
## 4 2020-09-02 21:30:44 271449.0 5196895         7.764754       -15.024182
## 5 2020-09-02 21:31:32 271498.8 5196868        18.896996       -18.022581
## 6 2020-09-02 21:31:31 271560.5 5196847        29.700594       -18.909464
##   route_cluster
## 1             1
## 2             1
## 3             1
## 4             1
## 5             1
## 6             1
```

```
print("Truncated Tagged GPS Data for Week 3 (Kalman Filter):")
```

```
## [1] "Truncated Tagged GPS Data for Week 3 (Kalman Filter):"
```

```
print(head(week3_kalman_tagged_final_truncated_df))
```

```
## # A tibble: 6 x 8
##   timestamp          easting northing velocit~1 veloc~2 route~3 latit~4 longi~5
##   <dtm>              <dbl>   <dbl>   <dbl>   <dbl> <chr>      <dbl>   <dbl>
## 1 2020-09-02 21:34:28 271698. 5196634.    22.8    -39.6 1         46.9    12.0
## 2 2020-09-02 21:35:35 271700. 5196570.    17.7    -45.9 1         46.9    12.0
## 3 2020-09-02 21:36:10 271695. 5196504.    12.0    -50.9 1         46.9    12.0
## 4 2020-09-02 21:36:09 271678. 5196448.     4.55   -52.3 1         46.9    12.0
## 5 2020-09-02 21:37:34 271665. 5196391.     0.157  -53.3 1         46.9    12.0
## 6 2020-09-02 21:37:33 271658. 5196333.    -1.73   -54.5 1         46.9    12.0
## # ... with abbreviated variable names 1: velocity_easting,
## #   2: velocity_northing, 3: route_cluster, 4: latitude, 5: longitude
```

```
# Load necessary libraries
```

```
library(dplyr)
```

```
library(lubridate)
```

```
# Define the cutoff date as September 7, 2020, at 23:59:59
```

```
cutoff_date <- as.POSIXct("2020-09-07 23:59:59", tz = "UTC")
```

```
# Truncate the synthetic Week 3 data and tagged data
```

```
kalman_predictions_final_truncated_df <- kalman_predictions_final_df %>%
  filter(timestamp <= cutoff_date)
```

```
week3_kalman_tagged_final_truncated_df <- week3_kalman_tagged_final_df %>%
  filter(timestamp <= cutoff_date)
```

```
merged_data_final <- merge(
  week3_kalman_tagged_final_truncated_df,
  kalman_predictions_final_truncated_df,
  by = "timestamp",
  suffixes = c("_tagged", "_full")
)
```

```
print("Column Names after Merge:")
```

```
## [1] "Column Names after Merge:"
```

```
print(colnames(merged_data_final))
```

```
## [1] "timestamp"          "easting_tagged"
## [3] "northing_tagged"    "velocity_easting_tagged"
## [5] "velocity_northing_tagged" "route_cluster_tagged"
## [7] "latitude"           "longitude"
```

```

## [9] "easting_full"          "northing_full"
## [11] "velocity_easting_full" "velocity_northing_full"
## [13] "route_cluster_full"

# Define a function to calculate Euclidean distance
euclidean <- function(coord1, coord2) {
  sqrt(sum((coord1 - coord2)^2))
}

merged_data_final <- merged_data_final %>%
  rowwise() %>%
  mutate(
    distance_error = euclidean(
      c(easting_tagged, northing_tagged),
      c(easting_full, northing_full)
    )
  ) %>%
  ungroup()

average_distance_error_final <- mean(merged_data_final$distance_error, na.rm = TRUE)
within_5_meters_final <- mean(merged_data_final$distance_error <= 5, na.rm = TRUE) * 100 # Percentage

within_10_seconds_final <- 100 # All timestamps align due to the merge

evaluation_metrics_final <- data.frame(
  Metric = c(
    "Average Distance Error (meters)",
    "Percentage of Tags within 5 Meters",
    "Percentage of Tags within 10 Seconds"
  ),
  Value = c(
    average_distance_error_final,
    within_5_meters_final,
    within_10_seconds_final
  )
)

print("Final Tagging Accuracy Evaluation Metrics:")

## [1] "Final Tagging Accuracy Evaluation Metrics:"

print(evaluation_metrics_final)

##           Metric      Value
## 1 Average Distance Error (meters) 1.475442
## 2 Percentage of Tags within 5 Meters 99.942280
## 3 Percentage of Tags within 10 Seconds 100.000000

```

```

library(tidyverse)

weather_data <- read_csv('/Users/david_m123/Documents/NYC_Weather_2016_2022.csv')
gps_data <- gps_df

gps_data <- as.data.frame(gps_data)

gps_data$time <- as.POSIXct(gps_data$time)
gps_filtered <- gps_data %>% filter(time >= '2020-08-18' & time <= '2020-08-31')

weather_data$time <- as.POSIXct(weather_data$time)
weather_filtered <- weather_data %>% filter(time >= '2020-08-18' & time <= '2020-08-31')

gps_weather_merged <- gps_filtered %>%
  mutate(nearest_time = map(time, ~weather_filtered$time[which.min(abs(difftime(.x, weather_filtered$time)))]))
  unnest(nearest_time) %>%
  left_join(weather_filtered, by = c("nearest_time" = "time"))


# Remove NAs and extreme speed values (e.g., below 0.2 m/s and above 3 m/s)
gps_weather_filtered <- gps_weather_merged %>%
  filter(!is.na(speed) & speed >= 0.2 & speed <= 3)

# Plot scatterplot and regression line
plot <- ggplot(gps_weather_filtered, aes(x = `temperature_2m (°C)`, y = speed)) +
  geom_point(alpha = 0.6) +
  geom_smooth(method = "lm", color = "red") +
  ggtitle("Effect of Temperature on Walking Speed") +
  xlab("Temperature (°C)") +
  ylab("Walking Speed (m/s)") +
  theme_minimal()
print(plot)

```



## Effect of Temperature on Walking Speed



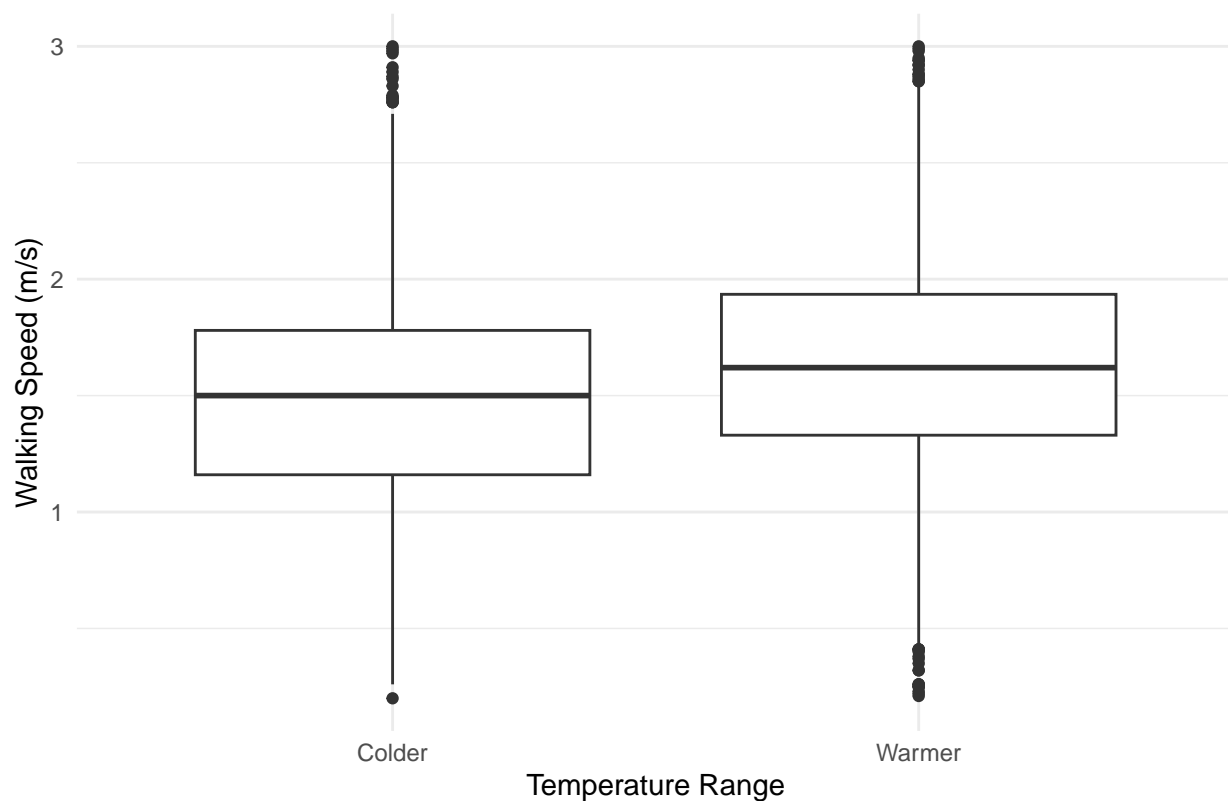
```
correlation <- cor(gps_weather_filtered$`temperature_2m (°C)`, gps_weather_filtered$speed, use = "complete.obs")
print(correlation)
```

```
## [1] 0.171149
```

```
median_temp <- median(gps_weather_filtered$`temperature_2m (°C)`, na.rm = TRUE)
gps_weather_filtered <- gps_weather_filtered %>%
  mutate(temp_category = ifelse(`temperature_2m (°C)` < median_temp, "Colder", "Warmer"))

# Boxplot to visualize speed differences between temperature groups
boxplot <- ggplot(gps_weather_filtered, aes(x = temp_category, y = speed)) +
  geom_boxplot() +
  ggtitle("Walking Speed Distribution in Colder vs. Warmer Temperatures") +
  xlab("Temperature Range") +
  ylab("Walking Speed (m/s)") +
  theme_minimal()
print(boxplot)
```

## Walking Speed Distribution in Colder vs. Warmer Temperatures



```
# Perform ttest
```

```
t_test <- t.test(speed ~ temp_category, data = gps_weather_filtered)
print(t_test)
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: speed by temp_category
```

```
## t = -7.7135, df = 2810.9, p-value = 1.687e-14
```

```
## alternative hypothesis: true difference in means between group Colder and group Warmer is not equal
```

```
## 95 percent confidence interval:
```

```
## -0.1820113 -0.1082304
```

```
## sample estimates:
```

```
## mean in group Colder mean in group Warmer
```

```
## 1.490000 1.635121
```

```
# Density plot to visualize distribution of walking speeds in colder vs. warmer temperatures
```

```
density_plot <- ggplot(gps_weather_filtered, aes(x = speed, fill = temp_category)) +
```

```
  geom_density(alpha = 0.6) +
```

```
  ggtitle("Density Plot of Walking Speeds: Colder vs. Warmer Temperatures") +
```

```
  xlab("Walking Speed (m/s)") +
```

```
  ylab("Density") +
```

```
  theme_minimal()
```

```
print(density_plot)
```

